

TASK-4:

Write MapReduce codes to perform the tasks using the files you've downloaded on your EMR Instance

Dataset used is: https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-05.csv

- 1) Get the dataset to with the following command:

```
wget https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-05.csv
```

- 2) Create python scripts by using the following command:

```
vim mrtask_a.py  
vim mrtask_b.py  
vim mrtask_c.py  
vim mrtask_d.py  
vim mrtask_e.py  
vim mrtask_f.py
```

- (a) mrtask_a.py

```
# Which vendors have the most trips, and what is the total revenue generated by  
that vendor?  
  
from mrjob.job import MRJob  
from mrjob.step import MRStep  
  
class MostTripsTotalRevenue(MRJob):  
  
    def steps(self):  
        return [  
            MRStep(mapper=self.mapper, reducer=self.reducer),  
            MRStep(reducer=self.final_reducer)  
        ]  
  
    def mapper(self, _, line):  
        if not line.startswith('VendorID'):  
            data = line.split(',')  
            vendor_id = data[0]  
            revenue = float(data[16])  
            yield vendor_id, revenue  
  
    def reducer(self, key, values):  
        yield None, (sum(values), key)  
  
    def final_reducer(self, _, values):  
        max_revenue, vendor_id = max(values)  
        yield vendor_id, max_revenue  
  
if __name__ == '__main__':  
    MostTripsTotalRevenue.run()
```

(b) mrtask_b.py

```
# Which pickup location generates the most revenue?

from mrjob.job import MRJob
from mrjob.step import MRStep

class MostRevenuePickupLocation(MRJob):

    def steps(self):
        return [
            MRStep(mapper=self.mapper, reducer=self.reducer),
            MRStep(reducer=self.final_reducer)
        ]

    def mapper(self, _, line):
        # Skip the header line
        if not line.startswith('VendorID'):
            fields = line.split(',')
            pickup_location = fields[7]
            revenue = float(fields[16])
            yield pickup_location, revenue

    def reducer(self, pickup_location, revenues):
        yield None, (sum(revenues), pickup_location)

    def final_reducer(self, _, max_revenues):
        max_revenue, pickup_location = max(max_revenues)
        yield pickup_location, max_revenue

if __name__ == '__main__':
    MostRevenuePickupLocation.run()
```

(c) mrtask_c.py

```
# What are the different payment types used by customers and their count? The
# final results should be in a sorted format.

from mrjob.job import MRJob
from mrjob.step import MRStep

class PaymentTypesCount(MRJob):

    def mapper(self, _, line):
        # Skip the header line
        if not line.startswith('VendorID'):
            fields = line.split(',')
            payment_type = fields[9]
            yield payment_type, 1

    def combiner(self, payment_type, counts):
        yield payment_type, sum(counts)

    def reducer(self, payment_type, counts):
        yield payment_type, sum(counts)

    def reducer_sort_results(self, payment_type, counts):
        yield None, (sum(counts), payment_type)

    def reducer_output_result(self, _, sorted_results):
        for count, payment_type in sorted(sorted_results, reverse=True):
            yield payment_type, count

    def steps(self):
        return [
            MRStep(mapper=self.mapper, combiner=self.combiner,
reducer=self.reducer),
            MRStep(reducer=self.reducer_sort_results),
            MRStep(reducer=self.reducer_output_result)
        ]

if __name__ == '__main__':
    PaymentTypesCount.run()
```

(d) mrtask_d.py

```
# What is the average trip time for different pickup locations?

from mrjob.job import MRJob
from datetime import datetime

class AverageTripTime(MRJob):

    def parse_datetime(self, datetime_str):
        formats = ['%d-%m-%Y %H:%M:%S', '%d-%m-%Y %H:%M', '%Y-%m-%d %H:%M', '%Y-%m-%d %H:%M:%S']
        for fmt in formats:
            try:
                return datetime.strptime(datetime_str, fmt)
            except ValueError:
                pass
        raise ValueError('no valid date format found')

    def mapper(self, _, line):
        # Skip the header line
        if not line.startswith('VendorID'):
            fields = line.split(',')
            pickup_location = fields[7]
            pickup_datetime = self.parse_datetime(fields[1])
            dropoff_datetime = self.parse_datetime(fields[2])
            trip_time = (dropoff_datetime - pickup_datetime).total_seconds() /
60.0

            yield pickup_location, (trip_time, 1)

    def combiner(self, pickup_location, trip_times):
        total_trip_time = 0
        total_count = 0
        for trip_time, count in trip_times:
            total_trip_time += trip_time
            total_count += count
        yield pickup_location, (total_trip_time, total_count)

    def reducer(self, pickup_location, trip_times):
        total_trip_time = 0
        total_count = 0
        for trip_time, count in trip_times:
            total_trip_time += trip_time
            total_count += count
        average_trip_time = total_trip_time / total_count
        yield pickup_location, average_trip_time

if __name__ == '__main__':
    AverageTripTime.run()
```

(e) mrtask_e.py

```
# Calculate the average tips to revenue ratio of the drivers for different pickup
locations in sorted format.

from mrjob.job import MRJob

class AverageTipsToRevenueRatio(MRJob):

    def mapper(self, _, line):
        # Skip the header line
        if not line.startswith('VendorID'):
            fields = line.split(',')
            pickup_location = fields[7]
            total_revenue = float(fields[16])
            tips = float(fields[13])
            yield pickup_location, (tips, total_revenue)

    def combiner(self, pickup_location, tips_revenues):
        total_tips = 0
        total_revenue = 0
```

```

    for tips, revenue in tips_revenues:
        total_tips += tips
        total_revenue += revenue
    yield pickup_location, (total_tips, total_revenue)

def reducer(self, pickup_location, tips_revenues):
    total_tips = 0
    total_revenue = 0
    for tips, revenue in tips_revenues:
        total_tips += tips
        total_revenue += revenue
    average_tips_to_revenue_ratio = total_tips / total_revenue
    yield pickup_location, average_tips_to_revenue_ratio

if __name__ == '__main__':
    AverageTipsToRevenueRatio.run()

```

(f) mrtask_f.py

```

# How does revenue vary over time? Calculate the average trip revenue per month -
# analysing it by hour of the day (day vs night) and the day of the week (weekday
# vs weekend).

from mrjob.job import MRJob
from datetime import datetime

class AverageRevenueOverTime(MRJob):

    def parse_datetime(self, datetime_str):
        formats = ['%d-%m-%Y %H:%M:%S', '%d-%m-%Y %H:%M', '%Y-%m-%d %H:%M', '%Y-
        %m-%d %H:%M:%S']
        for fmt in formats:
            try:
                return datetime.strptime(datetime_str, fmt)
            except ValueError:
                pass
        raise ValueError('no valid date format found')

    def mapper(self, _, line):
        # Skip the header line
        if not line.startswith('VendorID'):
            fields = line.split(',')
            revenue = float(fields[16])
            pickup_datetime = self.parse_datetime(fields[1])
            month = pickup_datetime.month
            hour = pickup_datetime.hour
            weekday = pickup_datetime.weekday()
            yield (month, hour, weekday), revenue

    def reducer(self, key, values):
        total_revenue = 0
        num_trips = 0

        for revenue in values:
            total_revenue += revenue
            num_trips += 1

        average_revenue = total_revenue / num_trips

        yield key, average_revenue

if __name__ == '__main__':
    AverageRevenueOverTime.run()

```

3) Before running these scripts we need to install mrjob by using the following commands:

```
sudo yum -y install python3-pip
sudo pip3 install mrjob

Updated:
python3-pip.noarch 0:20.2.2-1.amzn2.0.5

Complete!
[hadoop@ip-172-31-63-235 ~]$ sudo pip3 install mrjob
WARNING: Running pip install with root privileges is generally not a good idea. Try `pip3 install --user` instead.
Collecting mrjob
  Downloading mrjob-0.7.4-py2.py3-none-any.whl (439 kB)
    | 439 kB 21.8 MB/s
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib64/python3.7/site-packages (from mrjob) (5.3.1)
Installing collected packages: mrjob
  WARNING: The scripts mrjob, mrjob-3 and mrjob-3.7 are installed in '/usr/local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
Successfully installed mrjob-0.7.4
[hadoop@ip-172-31-63-235 ~]$
```

Questions:

(a) Which vendors have the most trips, and what is the total revenue generated by that vendor?

```
python mrtask_a.py yellow_tripdata_2017-05.csv > mrtask_a.txt

[hadoop@ip-172-31-63-235 ~]$ ls
mrtask_a.py mrtask_a.txt mrtask_b.py mrtask_c.py mrtask_d.py mrtask_e.py mrtask_f.py yellow_tripdata_2017-05.csv
[hadoop@ip-172-31-63-235 ~]$ python mrtask_a.py yellow_tripdata_2017-05.csv > mrtask_a.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/mrtask_a.hadoop.20240603.233826.200874
Running step 1 of 2...
Running step 2 of 2...
job output is in /tmp/mrtask_a.hadoop.20240603.233826.200874/output
Streaming final output from /tmp/mrtask_a.hadoop.20240603.233826.200874/output...
Removing temp directory /tmp/mrtask_a.hadoop.20240603.233826.200874...
[hadoop@ip-172-31-63-235 ~]$ cat mrtask_a.txt
"2"      92896777.54522054
[hadoop@ip-172-31-63-235 ~]$
```

Vendor "2" ie: VeriFone Inc. has the most trips and revenue generated is 92896777.545

(b) Which pickup location generates the most revenue?

```
python mrtask_b.py yellow_tripdata_2017-05.csv > mrtask_b.txt

[hadoop@ip-172-31-58-196 ~]$ python mrtask_b.py yellow_tripdata_2017-05.csv > mrtask_b.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/mrtask_b.hadoop.20240604.000704.853039
Running step 1 of 2...
Running step 2 of 2...
job output is in /tmp/mrtask_b.hadoop.20240604.000704.853039/output
Streaming final output from /tmp/mrtask_b.hadoop.20240604.000704.853039/output...
Removing temp directory /tmp/mrtask_b.hadoop.20240604.000704.853039...
[hadoop@ip-172-31-58-196 ~]$ cat mrtask_b.txt
"132"    14040591.220016211
[hadoop@ip-172-31-58-196 ~]$
```

Pickup location '132' generates the most revenue of 14040591.220016211

(c) What are the different payment types used by customers and their count? The final results should be in a sorted format.

```
python mrtask_c.py yellow_tripdata_2017-05.csv > mrtask_c.txt
[hadoop@ip-172-31-58-196 ~]$ python mrtask_c.py yellow_tripdata_2017-05.csv > mrtask_c.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/mrtask_c.hadoop.20240604.001225.060982
Running step 1 of 3...
Running step 2 of 3...
Running step 3 of 3...
job output is in /tmp/mrtask_c.hadoop.20240604.001225.060982/output
Streaming final output from /tmp/mrtask_c.hadoop.20240604.001225.060982/output...
Removing temp directory /tmp/mrtask_c.hadoop.20240604.001225.060982...
[hadoop@ip-172-31-58-196 ~]$ cat mrtask_c.txt
"1"      6780947
"2"      3250362
"3"      55027
"4"      15791
[hadoop@ip-172-31-58-196 ~]$ |
```

Different payment types used by customers with their count:

"1" = Credit Card	6780947
"2" = Cash	3250362
"3" = No Charge	55027
"4" = Dispute	15791

(d) What is the average trip time for different pickup locations?

```
python mrtask_d.py yellow_tripdata_2017-05.csv > mrtask_d.txt
```

```
[hadoop@ip-172-31-58-196 ~]$ cat mrtask_d.txt
"1"      12.681244881244885
"10"     56.4414970815586
"100"    16.301056898951547
"101"    10.8640350877193
"102"    15.480072463768114
"105"    22.727777777777778
"106"    14.046902507889056
"107"    14.914703800523975
"108"    17.747747747747745
"109"    13.047222222222222
"11"     11.318787878787878
"111"    15.347685185185185
"112"    15.342635286583413
"113"    15.655663780443364
"114"    16.265089572862763
"115"    9.849019607843136
"116"    15.213805066967888
"117"    9.020000000000001
"118"    16.645833333333336
"119"    14.638700564971753
"12"     25.11584025215252
"120"    12.665972222222223
"121"    16.845833333333333
"122"    11.954761904761904
"123"    12.698253968253967
"124"    35.719855072463766
"125"    16.582152788143137
"126"    32.175581395348836
"127"    17.629476052716164
"128"    15.57929292929293
"129"    14.331578947368426
"13"     20.914534001214356
"130"    37.98320691108302
"131"    11.690000000000003
"132"    46.805229740819875
"133"    17.532022849462365
"134"    16.221703296703296
"135"    15.368173258003768
"136"    9.693303571428572
"137"    14.412591446828932
"138"    41.79266819772832
"139"    18.8280512820513
"14"     18.910311918063318
"140"    14.544783434852864
"141"    12.887731943471524
"142"    14.724678434186814
"143"    13.87881130774534
```

"144"	17. 422410667367306
"145"	13. 919851029794044
"146"	15. 779948824343021
"147"	13. 003051643192489
"148"	16. 994242069527864
"149"	13. 134482758620688
"15"	14. 868589743589745
"150"	22. 067647058823532
"151"	13. 72286961525267
"152"	15. 034273318872025
"153"	12. 776791277258567
"154"	22. 930303030303033
"155"	23. 847222222222218
"156"	17. 777777777777778
"157"	24. 96510297482838
"158"	17. 056993321982915
"159"	12. 977884615384616
"16"	11. 23930817610063
"160"	17. 381399631675876
"161"	16. 501390496486135
"162"	15. 5755374379993
"163"	17. 007285407402943
"164"	16. 135913909640667
"165"	23. 146643518518516
"166"	16. 128181641420127
"167"	11. 731709401709402
"168"	11. 912370135587249
"169"	12. 008282828282828
"17"	14. 360986719350239
"170"	15. 23280819184737
"171"	11. 867179487179486
"172"	29. 295833333333333
"173"	12. 247398589065257
"174"	12. 091085271317832
"175"	11. 432758620689656
"176"	0. 6
"177"	20. 955774278215223
"178"	8. 291562500000001
"179"	15. 1499705882353
"18"	12. 653225806451612
"180"	27. 1835069444444446
"181"	15. 849603785417022
"182"	12. 2456
"183"	11. 927777777777779
"184"	7. 3466666666666666
"185"	15. 095833333333333
"186"	17. 3607280512866
"187"	1. 425
"188"	15. 907792044310169
"189"	16. 631431947577614
"19"	10. 768354430379746
"190"	26. 461985335589397
"191"	13. 486574074074074
"192"	13. 506842105263157
"193"	12. 075183464619935
"194"	27. 354013503375846
"195"	22. 571194503171245
"196"	18. 181564465408805
"197"	18. 036302681992343
"198"	13. 806195462478186
"2"	42. 2
"20"	17. 141091954022986
"200"	11. 855392156862747
"201"	0. 4937499999999997
"202"	15. 435855263157896
"203"	19. 28479532163743
"204"	0. 4799999999999987
"205"	18. 713008130081302
"206"	11. 648809523809524
"207"	8. 778993536472758
"208"	15. 355907172995781
"209"	20. 52999086340794
"21"	16. 197798742138364
"210"	18. 476602564102567
"211"	17. 805909853173578
"212"	24. 377777777777773
"213"	20. 40786163522013
"214"	8. 220833333333333
"215"	56. 68466165413534
"216"	34. 080833333333333
"217"	18. 35507590132827
"218"	17. 684177215189873
"219"	53. 02001915708813
"22"	22. 21085626911315
"220"	11. 665315315315317
"221"	19. 811904761904763
"222"	19. 394444444444442
"223"	17. 158481387478847
"224"	14. 423218777619141
"225"	15. 393991912189485
"226"	17. 062662041086035
"227"	15. 537459546925568
"228"	18. 118900966183574
"229"	13. 729634370452187
"23"	3. 605128205128205
"230"	18. 16775177926668
"231"	17. 842905702290793
"232"	17. 80712673224909
"233"	16. 130618672379082
"234"	15. 740532219114375
"235"	21. 51443736730361

```
"235" 21.514437736730361
"236" 13.552552898225786
"237" 12.943728784429947
"238" 13.635731925509624
"239" 13.766637412862702
"24" 14.086371798958675
"240" 16.91818181818182
"241" 15.555228758169934
"242" 13.335582822085891
"243" 17.764382566585954
"244" 17.97972746905131
"245" 23.330555555555552
"246" 16.162821706410472
"247" 16.876773493573694
"248" 13.558333333333332
"249" 15.335284107159861
"25" 17.17143186640833
"250" 14.828939393939397
"251" 2.1999999999999997
"252" 22.431190476190476
"253" 25.680555555555554
"254" 11.153439153439153
"255" 18.58691486586808
"256" 15.882166210462294
"257" 16.62408789386401
"258" 26.296310432569975
"259" 10.456666666666669
"26" 11.717295597484279
"260" 17.798491342362418
"261" 22.742345329876898
"262" 13.453498021610727
"263" 12.558972943698416
"264" 15.660701878308133
"265" 10.097814982973892
"27" 9.104166666666668
"28" 33.768
"29" 57.221774193548384
"3" 17.91923076923077
"30" 14.595833333333333
"31" 22.172916666666667
"32" 45.23925925925926
"33" 19.48258199014998
"34" 15.771012544802865
"35" 12.903932584269663
"36" 16.217423444449842
"37" 15.67785183038733
"38" 105.58416666666668
"39" 17.3040625
"4" 15.943539005698556
"40" 15.923511762519249
"41" 13.047298730655545
"42" 13.060410811865738
"43" 16.67567976383065
"44" 9.504166666666666
"45" 19.516391550098202
"46" 29.257407407407406
"47" 15.23495145631068
"48" 15.306237295695665
"49" 14.44308345010119
"5" 0.3
"50" 15.42537099704824
"51" 11.96068376068376
"52" 18.217088839941262
"53" 20.587239583333333
"54" 14.719335511982568
"55" 47.3187265917603
"56" 23.806329113924047
"57" 15.616666666666667
"58" 5.333333333333333
"59" 19.616666666666664
"6" 7.790277777777777
"60" 12.961559139784944
"61" 15.251830401496228
"62" 13.879747675962818
"63" 65.07303921568626
"64" 11.29107142857143
"65" 18.156792744967923
"66" 18.57747506019952
"67" 13.053431372549017
"68" 15.859693387283983
"69" 19.92417624521073
"7" 13.984743997175142
"70" 26.37702202480104
"71" 12.455172413793102
"72" 17.463793103448275
"73" 22.723684210526315
"74" 12.525644805036702
"75" 13.519075992588938
"76" 18.701871980676327
"77" 16.11434108527132
"78" 11.448753894080996
"79" 15.650715220659471
"8" 22.969314641744553
"80" 16.722509667955716
"81" 10.016666666666667
"82" 15.163276362823948
"83" 16.827437223042836
"84" 5.575
"85" 15.723731138545956
"86" 14.700000000000001
"87" 22.17575468660486
"88" 23.09566049232271
"89" 17.87714633298575
"9" 153.69927536231884
"90" 14.588625022901905
"91" 23.28269720101781
"92" 14.84195171026157
"93" 38.06967213114755
"94" 11.647380952380951
"95" 18.848467302452317
"96" 18.588888888888885
"97" 16.081618563620673
"98" 13.002500000000001
[hadoop@ip-172-31-58-196 ~]$
```


(e) Calculate the average tips to revenue ratio of the drivers for different pickup locations in sorted format.

```
python mrtask_e.py yellow_tripdata_2017-05.csv > mrtask_e.txt
```

```
[hadoop@ip-172-31-58-196 ~]$ ls
mrtask_a.py mrtask_b.txt mrtask_c.txt mrtask_d.txt mrtask_f.py          yellow_tripdata_2017-05.csv.1
mrtask_b.py mrtask_c.py mrtask_d.py mrtask_e.py yellow_tripdata_2017-05.csv
[hadoop@ip-172-31-58-196 ~]$ python mrtask_e.py yellow_tripdata_2017-05.csv > mrtask_e.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/mrtask_e.hadoop.20240604.011220.382512
Running step 1 of 1...
job output is in /tmp/mrtask_e.hadoop.20240604.011220.382512/output
Streaming final output from /tmp/mrtask_e.hadoop.20240604.011220.382512/output...
Removing temp directory /tmp/mrtask_e.hadoop.20240604.011220.382512...
[hadoop@ip-172-31-58-196 ~]$ |
[hadoop@ip-172-31-58-196 ~]$ cat mrtask_e.txt
"1"      0.11961918264336044
"10"     0.1031314953272488
"100"    0.10054811858030266
"101"    0.1538484771604267
"102"    0.09686867923894887
"105"    0.07175925925925926
"106"    0.11203391656570175
"107"    0.11929564967810695
"108"    0.12120884627098828
"109"    0.23737785016286642
"11"     0.058159884648828326
"111"    0.077659754354142
"112"    0.10900187462682069
"113"    0.11769752457713911
"114"    0.11571668143567702
"115"    0.1014877461571702
"116"    0.09140144834181262
"117"    0.029865191166550988
"118"    0.03324247058041353
"119"    0.06439098561727422
"12"     0.09033344274559964
"120"    0.14485426649668087
"121"    0.06025386778698323
"122"    0.035274937009041045
"123"    0.0750228876958997
"124"    0.09886020334918795
"125"    0.1223467893609015
"126"    0.05493403420192623
"127"    0.08833111068775
"128"    0.10511320192841678
"129"    0.06647065796213625
"13"     0.12390141102889671
"130"    0.1073558396219219
"131"    0.10384807204137783
"132"    0.1025660513460536
"133"    0.09147281347725217
"134"    0.10314438578934823
"135"    0.0785122214655781
"136"    0.032600578121018825
"137"    0.11253077758026406
"138"    0.1304834446631976
"139"    0.025278845473624637
```

```

"140" 0.1110704470853594
"141" 0.11186974122790275
"142" 0.1131558975317282
"143" 0.11473162975887714
"144" 0.1105272721674073
"145" 0.09866671202718473
"146" 0.09229805528093973
"147" 0.04671430427439263
"148" 0.11347366987186253
"149" 0.12259087582337157
"15" 0.12542820971105154
"150" 0.10633375877472878
"151" 0.10794854174655912
"152" 0.08497213797829244
"153" 0.08409909578769388
"154" 0.07400473993394348
"155" 0.09130948944530835
"156" 0.07496251874062969
"157" 0.11154334990834797
"158" 0.1177428105347338
"159" 0.04880437380957894
"16" 0.1296825635625094
"160" 0.09369565599255213
"161" 0.09989106877197958
"162" 0.11962265897953572
"163" 0.11101476605573193
"164" 0.11138449806107008
"165" 0.07196634804086764
"166" 0.11344656333427505
"167" 0.0356042774025466
"168" 0.050410277807923884
"169" 0.11765956298165817
"17" 0.08304806476027186
"170" 0.11737627919224174
"171" 0.04416689447785676
"172" 0.27196083763937995
"173" 0.0662888212150767
"174" 0.09292588232864547
"175" 0.08816570904350345
"176" 0.19982277359326542
"177" 0.0985123820057963
"178" 0.03941897087008663
"179" 0.08647112167836499
"18" 0.07413049841666887
"180" 0.10306335321889983
"181" 0.11383678943060177
"182" 0.04935382990636337
"183" 0.04058310746465461
"184" 0.03374233128834356
"185" 0.07860584480628201
"186" 0.10968385132719934
"187" 0.03211303789338471
"188" 0.07172143661981017
"189" 0.10932535246060386
"19" 0.11032972645756875
"190" 0.10117750519544694
"191" 0.05438239409744308
"192" 0.0720589495330306
"193" 0.03922346473669742
"194" 0.12076731965739748
"195" 0.09492841545380133
"196" 0.09427848086047703
"197" 0.10215364731577099
"198" 0.10078643815066789
"2" 0.11989767545323099
"20" 0.0736013986013986
"200" 0.13187606166264348
"201" 0.132044713553796
"202" 0.09899140045991586
"203" 0.048087416271543615
"204" 0.0684931506849315
"205" 0.07523246599344224
"206" 0.08833170762908359
"207" 0.0815078125388226
"208" 0.08291166272142231
"209" 0.11297948250644907
"21" 0.06386360722984234
"210" 0.07934380776340108
"211" 0.11310320504679515
"212" 0.06409624702307627
"213" 0.11780938978546256
"214" 0.02993836219547989
"215" 0.10353565292295142
"216" 0.11008942973165864
"217" 0.0786879019937758
"218" 0.19235819846724786
"219" 0.1049626881741011
"22" 0.05041866234795492
"220" 0.056836272960188075
"221" 0.10844779603195978
"222" 0.11278732656864776
"223" 0.09822645175828154
"224" 0.11425442486544808
"225" 0.08099702729092222
"226" 0.08354577861128779
"227" 0.08721017303002598
"228" 0.09978962780363482
"229" 0.11203502738122283
"23" 0.08146106133700898
"230" 0.10662429938743652
"231" 0.11828076056384894

```

```

"232" 0.10577578661446632
"233" 0.11514074932415078
"234" 0.12096822968790891
"235" 0.0557415553911836
"236" 0.11249379678942588
"237" 0.1096331664258742
"238" 0.10834063914173245
"239" 0.11534544071014695
"24" 0.10648126258555479
"240" 0.10447949050358085
"241" 0.06993302518630319
"242" 0.07890211996979177
"243" 0.09530565669975881
"244" 0.09546944114505117
"245" 0.0
"246" 0.11995315301676546
"247" 0.07335984971767694
"248" 0.052623959464350334
"249" 0.11923576897943587
"25" 0.11052464311894397
"250" 0.043282871250452847
"251" 0.0019544811624019544
"252" 0.08558204423759197
"253" 0.06711331976762965
"254" 0.0972480119905075
"255" 0.11601004961517414
"256" 0.1093043993422745
"257" 0.10482648999719887
"258" 0.15786461137400468
"259" 0.12318663572582067
"26" 0.05332139249601685
"260" 0.07205229212930792
"261" 0.10801526749738075
"262" 0.11380613560114887
"263" 0.1142608495023124
"264" 0.11230316290901997
"265" 0.10968911731785466
"27" 0.07044812927929389
"28" 0.10315928607679181
"29" 0.04496955778040933
"3" 0.0352269585440825
"30" 0.431189363995688
"31" 0.09241457441314165
"32" 0.07388668844623049
"33" 0.1204236725267228
"34" 0.11469105842822479
"35" 0.17581828771112398
"36" 0.10412970427360184
"37" 0.103438070866592
"38" 0.03832478224555542
"39" 0.11607313499394789

"39" 0.11607313499394789
"4" 0.10621226787883002
"40" 0.11953106379554157
"41" 0.09274584418153507
"42" 0.07096965920546744
"43" 0.10407835112043434
"44" 0.07601977750309023
"45" 0.0983846147461719
"46" 0.12534064929126656
"47" 0.03340712045748017
"48" 0.10582758808855866
"49" 0.09681013151351817
"5" 0.711159737417943
"50" 0.10923919835836489
"51" 0.054687879641958155
"52" 0.13055804576734864
"53" 0.09939808622286242
"54" 0.11968717470040443
"55" 0.10775881095465412
"56" 0.09187656958656626
"57" 0.0982393882217696
"58" 0.023456558453743663
"59" 0.022958827634333566
"6" 0.04668850781897014
"60" 0.028830566964685292
"61" 0.09614549547751258
"62" 0.07257887627997188
"63" 0.07430008293013912
"64" 0.07777953179990528
"65" 0.1064885388089956
"66" 0.11474873960246652
"67" 0.08016136157322679
"68" 0.1143179862016906
"69" 0.03888023995673477
"7" 0.08055330999297121
"70" 0.10558719618955761
"71" 0.057126398156593525
"72" 0.12979482386561036
"73" 0.06660273315751619
"74" 0.07309332094094213
"75" 0.09233854474158283
"76" 0.15699972823625333
"77" 0.03758990781571329
"78" 0.03340662069771519
"79" 0.11726023735794265
"8" 0.1128809078373831
"80" 0.11230199593086634
"81" 0.1314140628199353
"82" 0.0640252588354194

```

```
"83" 0.060647576771211555
"84" 0.0006641576267434137
"85" 0.0568739741017691
"86" 0.03836325514349365
"87" 0.12371478220911764
"88" 0.11535120965220538
"89" 0.08902658805206363
"9" 0.07875511622519114
"90" 0.11768850776024835
"91" 0.10308391768201007
"92" 0.056615381413745704
"93" 0.1071444344921823
"94" 0.038084346952212626
"95" 0.07928530356245143
"96" 0.09975598325818912
"97" 0.09868461512594295
"98" 0.06067388225420083
[hadoop@ip-172-31-58-196 ~]$ |
```

(f) How does revenue vary over time? Calculate the average trip revenue per month - analysing it by hour of the day (day vs night) and the day of the week (weekday vs weekend).

```
python mrtask_f.py yellow_tripdata_2017-05.csv > mrtask_f.txt
```

```
[hadoop@ip-172-31-58-196 ~]$ python mrtask_f.py yellow_tripdata_2017-05.csv > mrtask_f.txt
No configs found; falling back on auto-configuration
No configs specified for inline runner
Creating temp directory /tmp/mrtask_f.hadoop.20240604.013608.691540
Running step 1 of 1...
job output is in /tmp/mrtask_f.hadoop.20240604.013608.691540/output
Streaming final output from /tmp/mrtask_f.hadoop.20240604.013608.691540/output...
Removing temp directory /tmp/mrtask_f.hadoop.20240604.013608.691540...
[hadoop@ip-172-31-58-196 ~]$ cat mrtask_f.txt
[5, 0, 0] 19.09149097120672
[5, 0, 1] 19.42010595117413
[5, 0, 2] 17.686158986637274
[5, 0, 3] 17.25599944664103
[5, 0, 4] 18.099999027301898
[5, 0, 5] 17.105128026116333
[5, 0, 6] 15.658808411947584
[5, 1, 0] 17.87588900705367
[5, 1, 1] 19.429909785928686
[5, 1, 2] 17.373021160869822
[5, 1, 3] 16.436848288253056
[5, 1, 4] 18.217479700823546
[5, 1, 5] 16.275959370214274
[5, 1, 6] 15.141293067405417
[5, 10, 0] 16.241741228653577
[5, 10, 1] 16.639627234011154
[5, 10, 2] 27.168365753878387
[5, 10, 3] 16.499299606770162
[5, 10, 4] 16.48680345506514
[5, 10, 5] 14.029692293163937
[5, 10, 6] 14.908801744148995
[5, 11, 0] 16.680433054406084
[5, 11, 1] 17.19449236461242
[5, 11, 2] 16.92334503331534
[5, 11, 3] 18.596232180453956
[5, 11, 4] 16.728098769455887
[5, 11, 5] 14.021549507826318
[5, 11, 6] 14.98238805727666
[5, 12, 0] 16.480360929345835
[5, 12, 1] 17.22963705731718
```

```
[5, 12, 2] 17.05563095494162
[5, 12, 3] 17.975730053716557
[5, 12, 4] 17.69103689287293
[5, 12, 5] 14.499792055020734
[5, 12, 6] 15.27335388620518
[5, 13, 0] 16.795085865627826
[5, 13, 1] 17.27948235089487
[5, 13, 2] 17.774827762048346
[5, 13, 3] 18.595673825475238
[5, 13, 4] 18.205408289793773
[5, 13, 5] 14.99034880898968
[5, 13, 6] 15.937601258660894
[5, 14, 0] 17.20858440343585
[5, 14, 1] 17.34785479975938
[5, 14, 2] 18.110961353384074
[5, 14, 3] 19.033494081619924
[5, 14, 4] 18.613387909146194
[5, 14, 5] 15.336584485461453
[5, 14, 6] 17.128896489169726
[5, 15, 0] 16.91500182066798
[5, 15, 1] 17.307682451791745
[5, 15, 2] 17.75765307206797
[5, 15, 3] 18.750789042356082
[5, 15, 4] 18.479680845892585
[5, 15, 5] 15.509464318824842
[5, 15, 6] 17.056592424577616
[5, 16, 0] 18.26141013123625
[5, 16, 1] 18.859940648116382
[5, 16, 2] 19.6542708916958
[5, 16, 3] 20.04885999577759
[5, 16, 4] 19.834795454141705
[5, 16, 5] 15.635747754243198
[5, 16, 6] 17.32405836477559
[5, 17, 0] 17.555959492507064
[5, 17, 1] 17.747241962602008
[5, 17, 2] 18.3960718339617
[5, 17, 3] 18.887732509348027
[5, 17, 4] 18.411963137349552
[5, 17, 5] 15.3593093037344426
[5, 17, 6] 17.4439645341407
[5, 18, 0] 16.500203407324047
[5, 18, 1] 16.477160631363912
[5, 18, 2] 17.053779035240108
[5, 18, 3] 17.221604263570683
[5, 18, 4] 16.83762356868249
[5, 18, 5] 14.73267798036185
[5, 18, 6] 16.363879207901295
[5, 19, 0] 16.117071206744605
[5, 19, 1] 15.933338279383616
[5, 19, 2] 16.270620125859278
[5, 19, 3] 16.472398718229226
```

```
[5, 19, 4] 16.161423076937226
[5, 19, 5] 14.486719040663449
[5, 19, 6] 16.246780024174765
[5, 2, 0] 17.069302560401354
[5, 2, 1] 17.102999568033265
[5, 2, 2] 16.490724900576904
[5, 2, 3] 15.885527411273001
[5, 2, 4] 17.510186185909916
[5, 2, 5] 15.787822292178655
[5, 2, 6] 15.219828879150024
[5, 20, 0] 16.278329889820224
[5, 20, 1] 15.92229313006401
[5, 20, 2] 15.999283947134947
[5, 20, 3] 16.363236084025758
[5, 20, 4] 16.041364614564298
[5, 20, 5] 15.008297507476271
[5, 20, 6] 16.514845325711754
[5, 21, 0] 16.81774921469818
[5, 21, 1] 16.388015828166143
[5, 21, 2] 16.231406528150977
[5, 21, 3] 17.121041382564638
[5, 21, 4] 16.579383172574527
[5, 21, 5] 15.030104427307885
[5, 21, 6] 16.83408532267027
[5, 22, 0] 17.337623268018195
[5, 22, 1] 17.00356693816382
[5, 22, 2] 16.730750002516206
[5, 22, 3] 17.699894119517573
[5, 22, 4] 16.72930392121595
[5, 22, 5] 15.161717603507384
[5, 22, 6] 17.47747548027094
[5, 23, 0] 18.734166782523527
[5, 23, 1] 17.364087663335088
[5, 23, 2] 17.252883134436264
[5, 23, 3] 18.14697136778204
[5, 23, 4] 17.248078089078962
[5, 23, 5] 15.846046419400455
[5, 23, 6] 18.86490867580518
[5, 3, 0] 17.010436833469626
[5, 3, 1] 17.389016241913282
[5, 3, 2] 16.62262594124502
[5, 3, 3] 16.37661024166937
[5, 3, 4] 17.67973900120963
[5, 3, 5] 16.147548793518364
[5, 3, 6] 15.822056911276535
[5, 4, 0] 21.292999725297417
[5, 4, 1] 20.068907006970854
[5, 4, 2] 19.732507288628188
[5, 4, 3] 20.829968819597944
[5, 4, 4] 20.9156692772681466
[5, 4, 5] 17.825446810289648
```

```
[5, 4, 6] 17.38188976038157
[5, 5, 0] 21.593018355659638
[5, 5, 1] 18.81002788103778
[5, 5, 2] 18.894916666663246
[5, 5, 3] 19.845702192891927
[5, 5, 4] 21.227005213121405
[5, 5, 5] 20.992848200309997
[5, 5, 6] 21.125313003037945
[5, 6, 0] 16.36666989228269
[5, 6, 1] 14.978345859721552
[5, 6, 2] 14.767823020740611
[5, 6, 3] 15.288311117881104
[5, 6, 4] 16.389897123112565
[5, 6, 5] 19.32208822201787
[5, 6, 6] 19.992739240104342
[5, 7, 0] 15.307352827781946
[5, 7, 1] 14.68863340021126
[5, 7, 2] 14.46050088621336
[5, 7, 3] 14.619068724817286
[5, 7, 4] 14.939835910422332
[5, 7, 5] 16.300197449100637
[5, 7, 6] 18.066767322725
[5, 8, 0] 15.534355540424196
[5, 8, 1] 15.310034719953286
[5, 8, 2] 14.991955112543714
[5, 8, 3] 15.095595443793671
[5, 8, 4] 15.144876788509523
[5, 8, 5] 14.67262836614724
[5, 8, 6] 15.685100979728976
[5, 9, 0] 15.79006959000546
[5, 9, 1] 15.809253019981295
[5, 9, 2] 15.622612773623201
[5, 9, 3] 15.764076749004406
[5, 9, 4] 15.912008246957104
[5, 9, 5] 13.814643752241684
[5, 9, 6] 14.779713618189996
[hadoop@ip-172-31-58-196 ~]$ |
```